

Announcements

Reminder: ps3 due Thursday 10/8 at midnight (Boston)

Assignment Project Exam Help

- ps4 out Thursday, due 10/15 (1 week)
- Lab this week – neural network learning
- ps3 self-grading form out Monday, due 10/19

<https://powcoder.com>

Add WeChat powcoder

A stylized illustration of a neural network. It features a central neuron with a purple nucleus and a blue cytoplasm, surrounded by a dense web of blue dendrites and axons. Several bright orange sparks or light points are scattered throughout the network, suggesting active connections or data flow.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Neural Networks III

Today: Outline

- **Neural networks cont'd**

Assignment Project Exam Help

- **Types of networks:** Feed-forward networks, convolutional networks, recurrent networks

<https://powcoder.com>

Add WeChat powcoder

- **ConvNets:** multiplication vs convolution; filters (or kernels); convolutional layers; 1D and 2D convolution; pooling layers; LeNet, CIFAR10Net



Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Neural Networks III

Network Architectures

Neural networks: recap




Learn parameters via gradient descent

$$\min_{\Theta} J(\Theta)$$

Assignment Project Exam Help
<https://powcoder.com>

Add WeChat powcoder

Backpropagation efficiently computes cost (forward pass) and gradient (backward pass)




$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$$


Network architectures

Feed-forward


Fully connected



Recurrent



Convolutional





Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Neural Networks III

Convolutional Architectures

Multiplication vs convolution

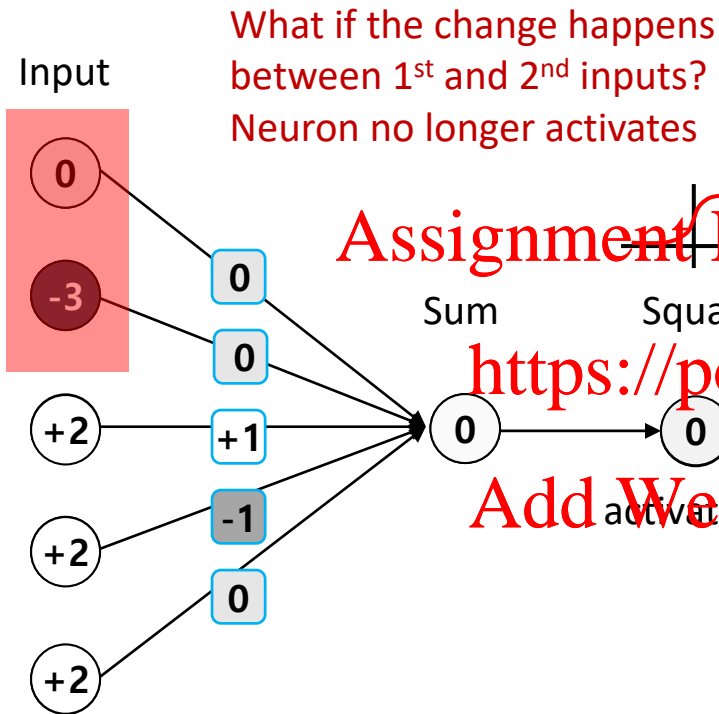
Input



- Recall, a neuron can be thought of as learning to spot certain features in the input
- E.g., this neuron detects change from high to low (light to dark) between 3rd and 4th inputs

<https://powcoder.com>
Add WeChat powcoder

Multiplication vs convolution



- Must have a new neuron for each new location of pattern???

- This is not efficient

- Solution: use convolution instead of multiplication

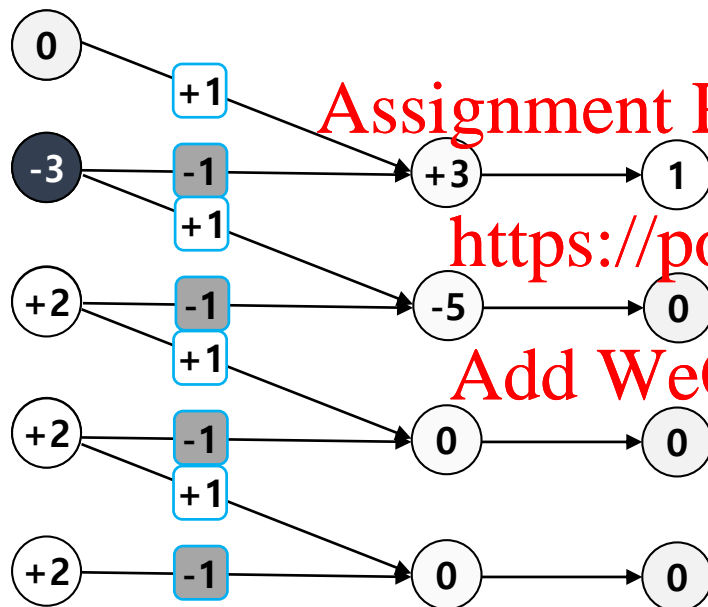
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

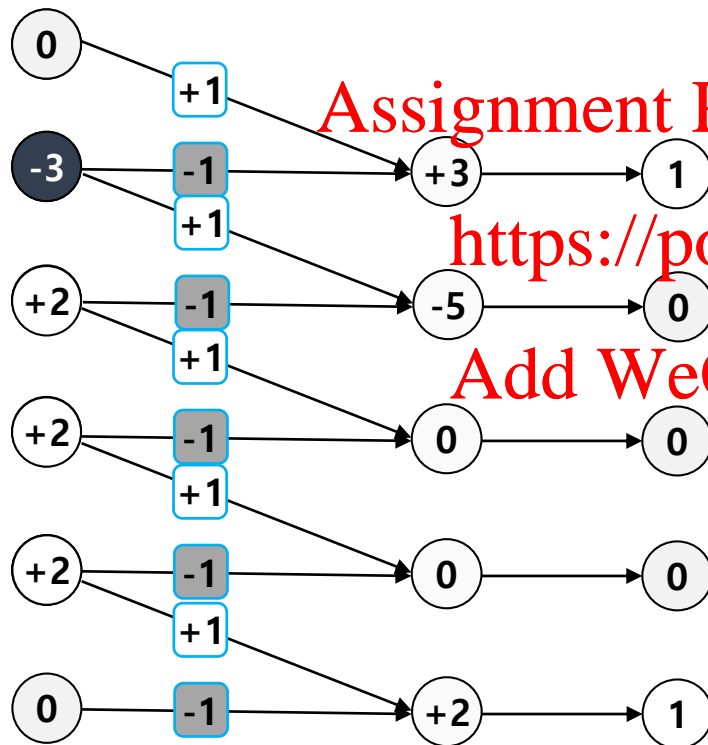
Input



- New weights are of size 2 x 1; called **filter**, or **kernel**
- New output is the size of input minus 1 because of **boundary**
- New convolutional neurons all share the same weights! This is much more efficient; we learn the weights once instead of many times for each position

Multiplication vs convolution

Padded
Input



- New output is the size of input minus 1 because of **boundary**
- We can fix the boundary effect by padding the input with 0 and adding one more neuron

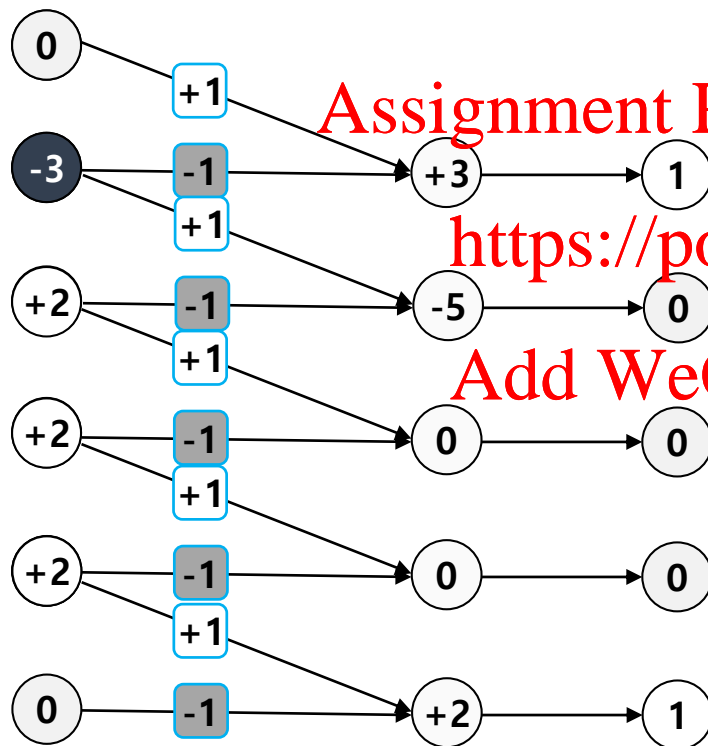
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**




Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**
- Stride can be larger, e.g. here is stride 2


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

Padded
Input




- We can add another filter, this time to detect opposite change with weights $[-1 \ +1]$
- Unique filters are called **channels**

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution

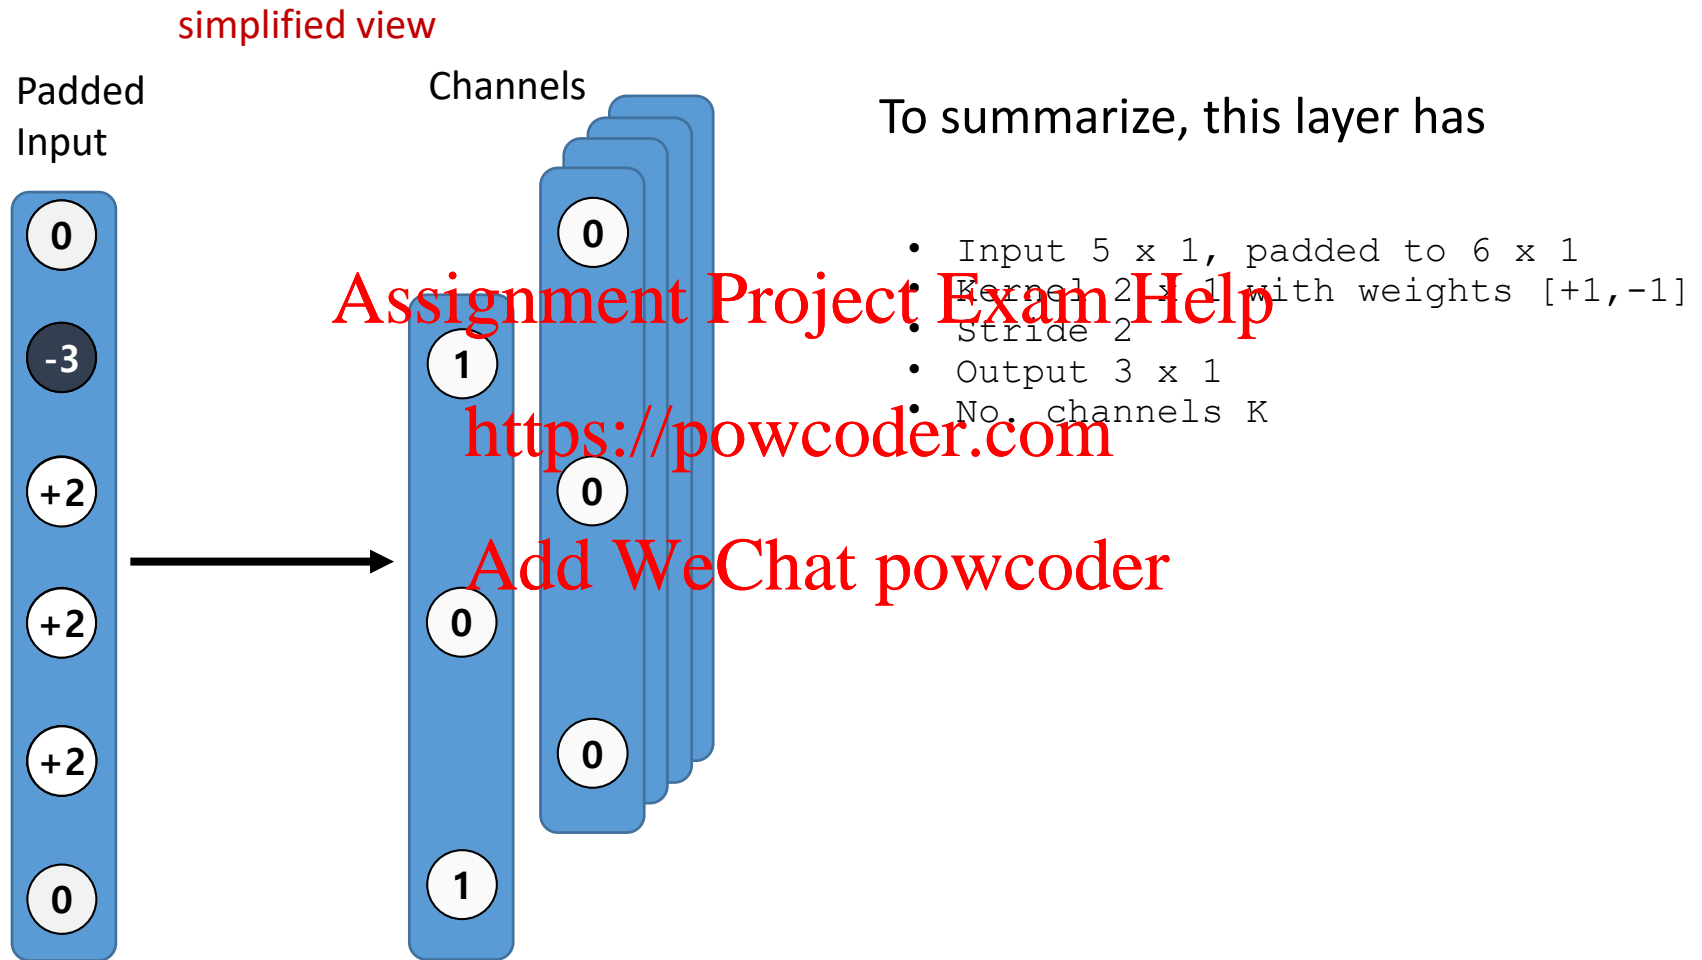


- We can add another filter, this time to detect opposite change with weights $[-1 \ +1]$
- Unique filters are called **channels**

<https://powcoder.com>

Add WeChat powcoder

Multiplication vs convolution





Assignment Project Exam Help

<https://powcoder.com>

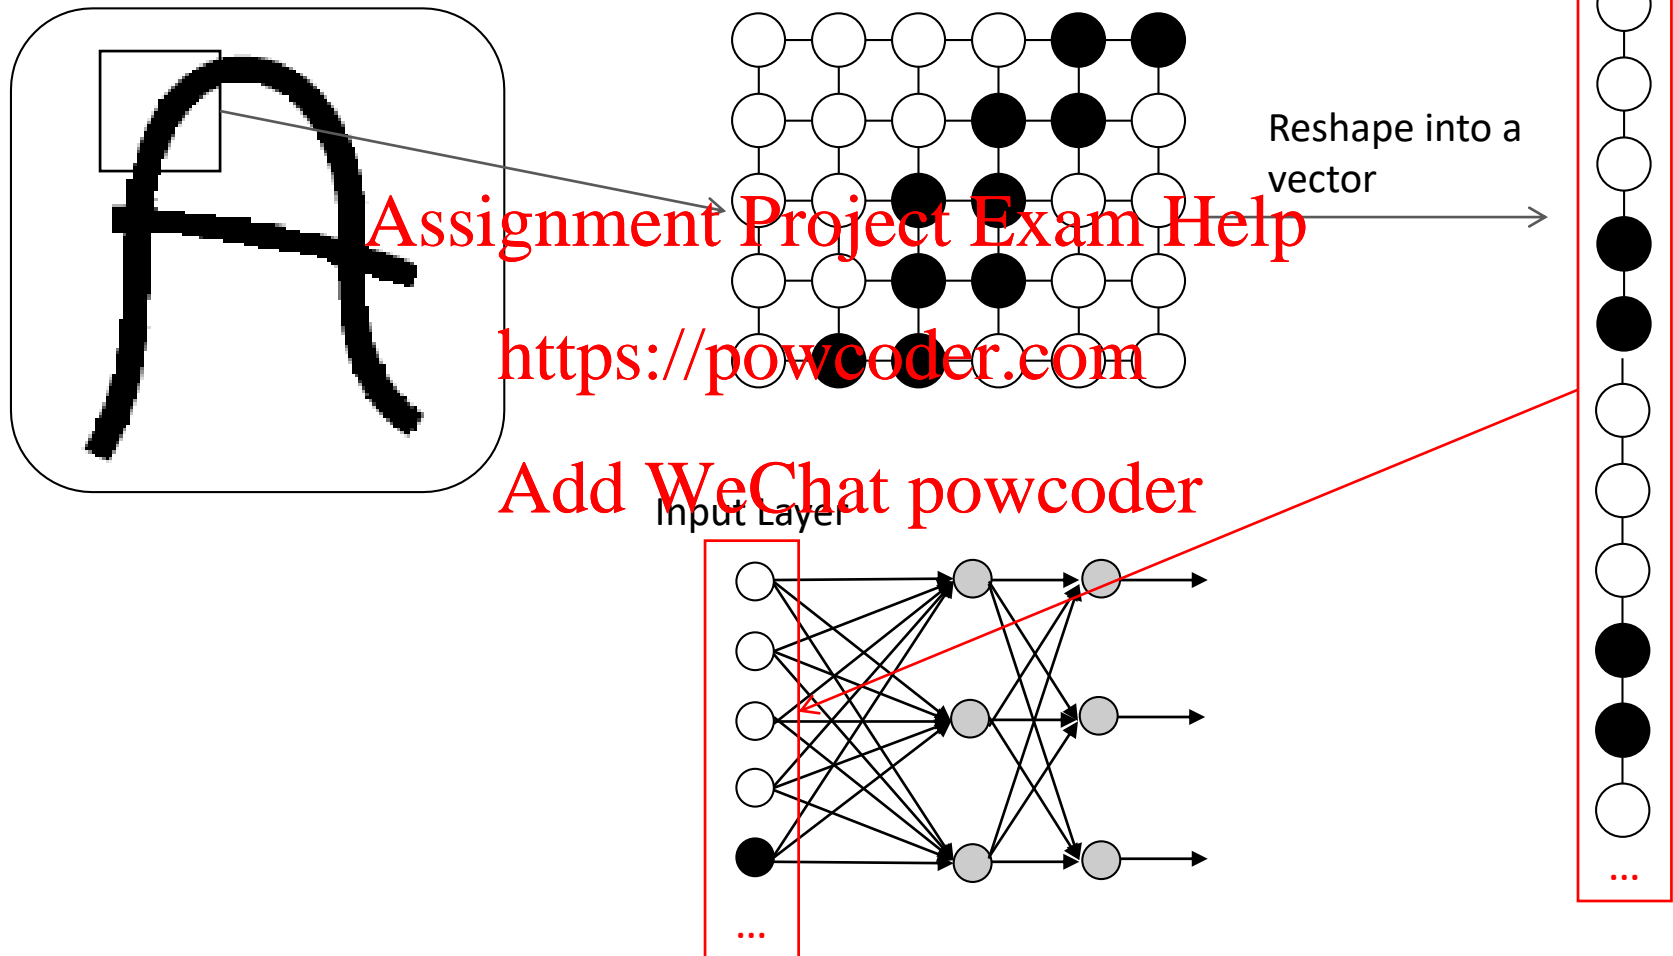
Add WeChat powcoder

Convolutional Neural Networks

For images and other 2-D signals


Representing images

Fully connected



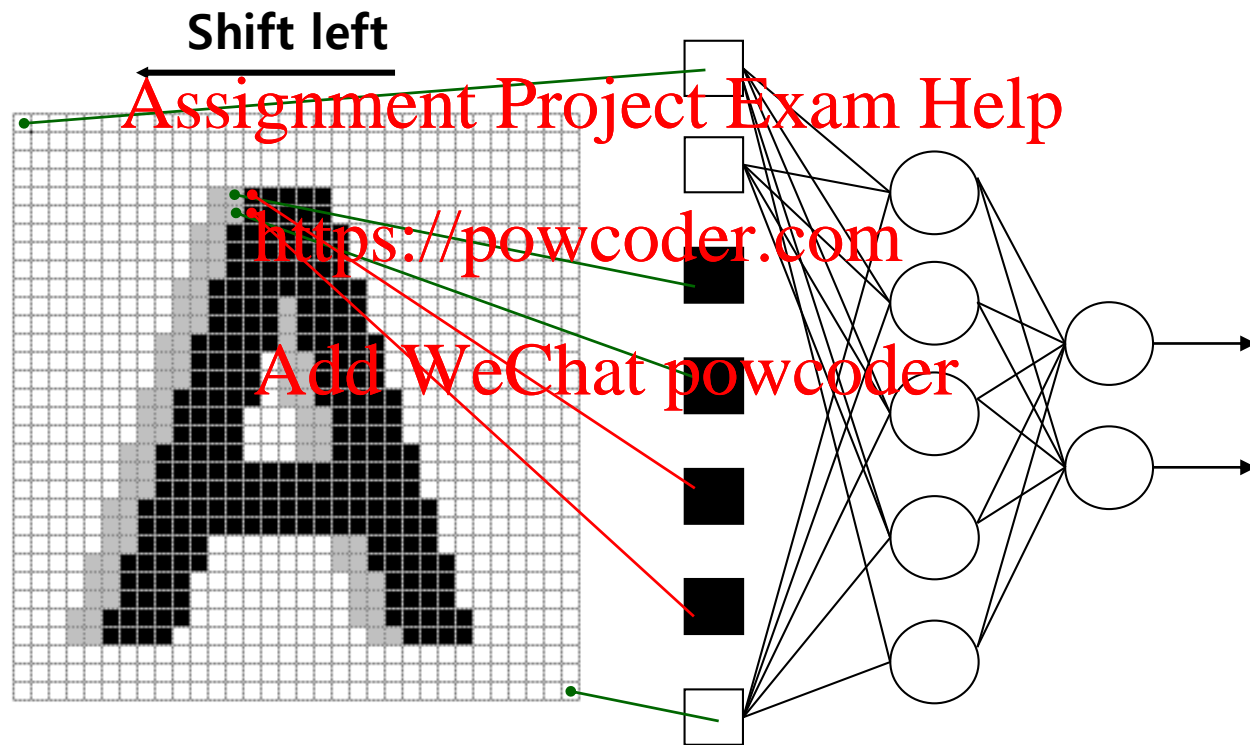
2D Input: fully connected network

Vectorize input by copying rows into a single column

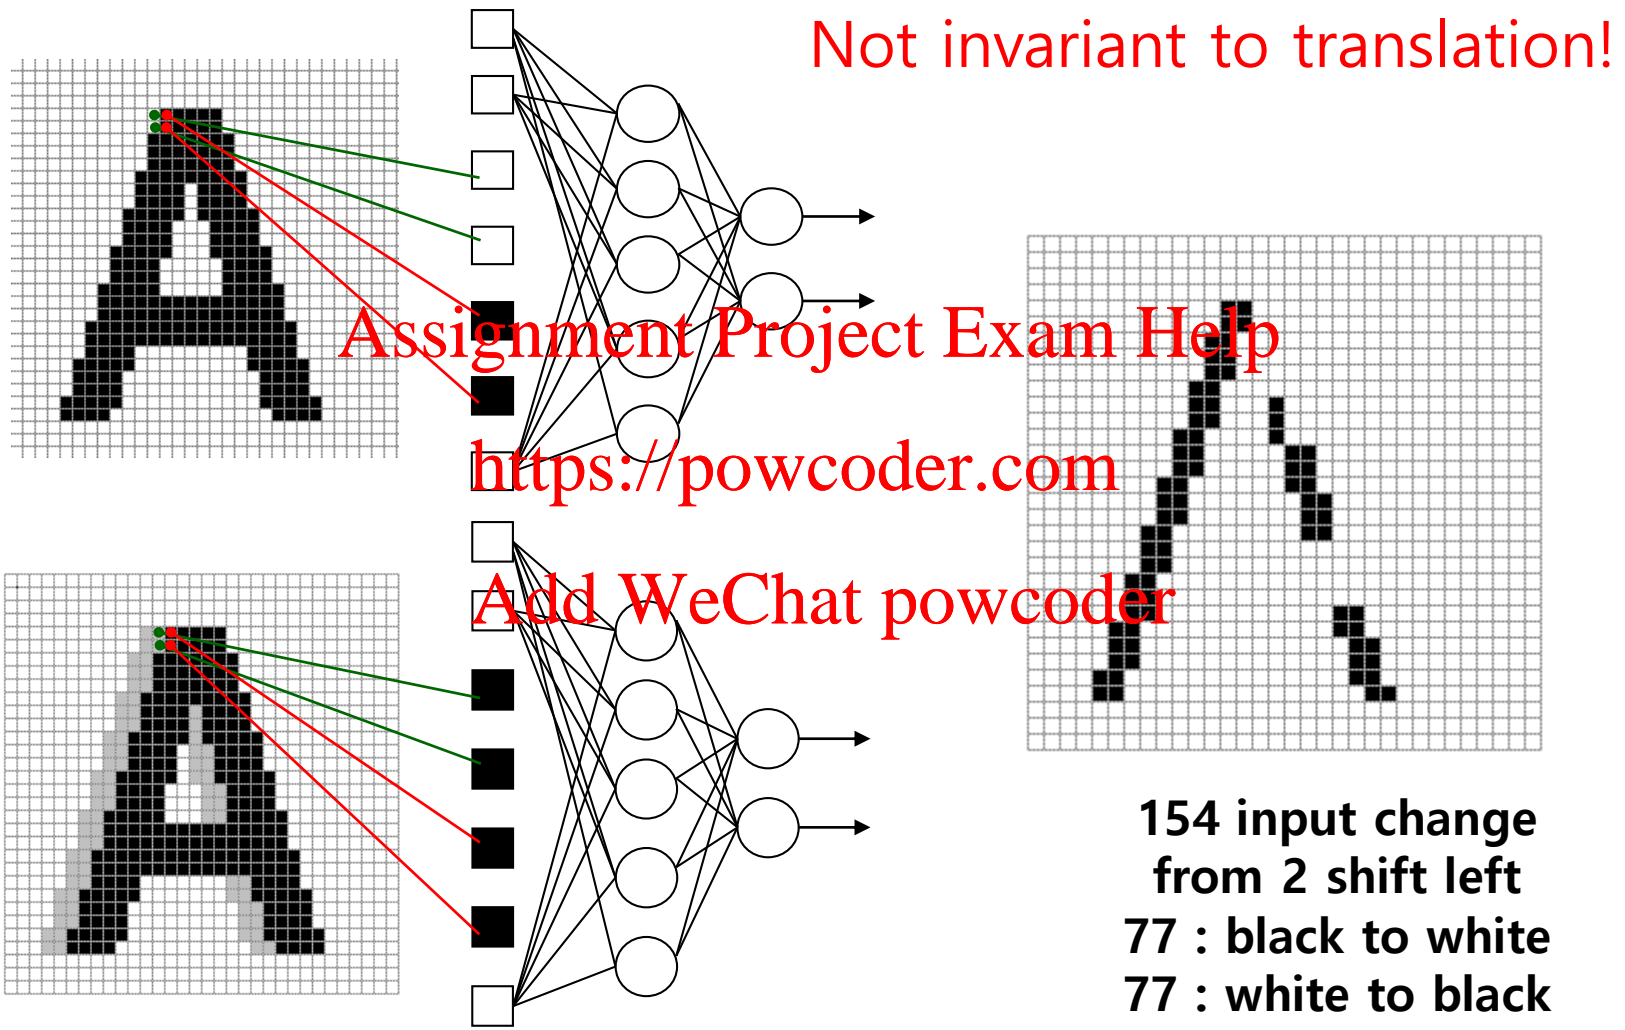


2D Input: fully connected network

Problem: shifting, scaling, and other distortion changes location of features



2D Input: fully connected network



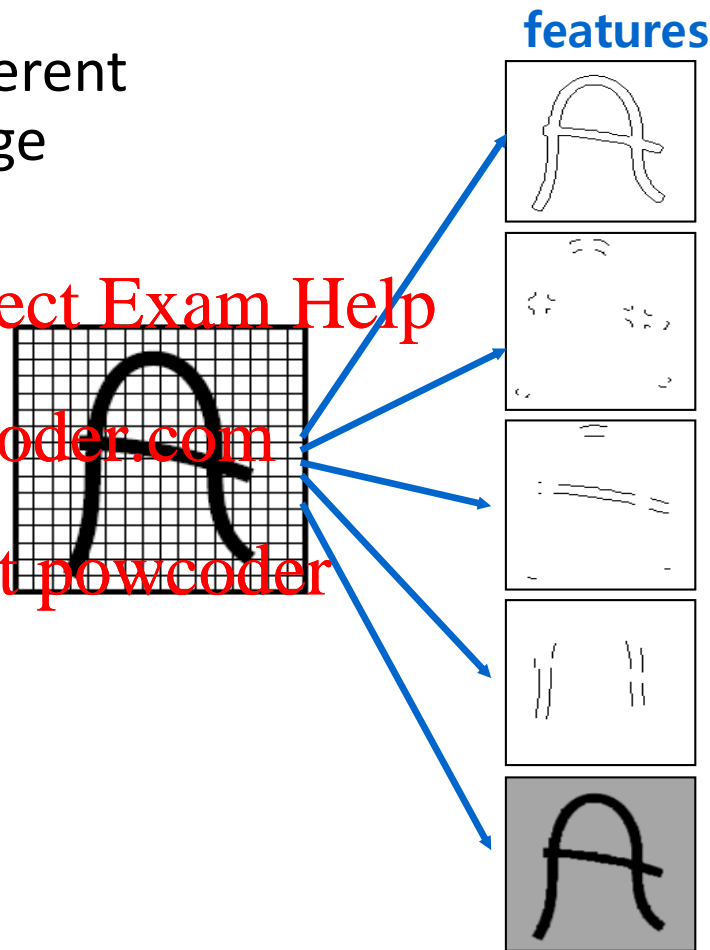
Convolution layer in 2D

- detect the same feature at different positions in the input, e.g. image
- preserve input topology


Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



Convolution layer in 2D



Input image




Output map

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Convolution layer in 2D



Input image




Output map

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Output looks like
an image

$$a = f(w_{11}x_{11} + w_{12}x_{12} + w_{13}x_{13} + \cdots w_{33}x_{33})$$

What weights correspond to these output maps?

*These are output maps
before thresholding*

*Hint: filters look like the
input they fire on*





Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


$$\frac{\partial f}{\partial x}(x, y)$$

∂x



$$\frac{\partial f}{\partial y}(x, y)$$

∂y



What will the output map look like?



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



filter

Input

What will the output map look like?

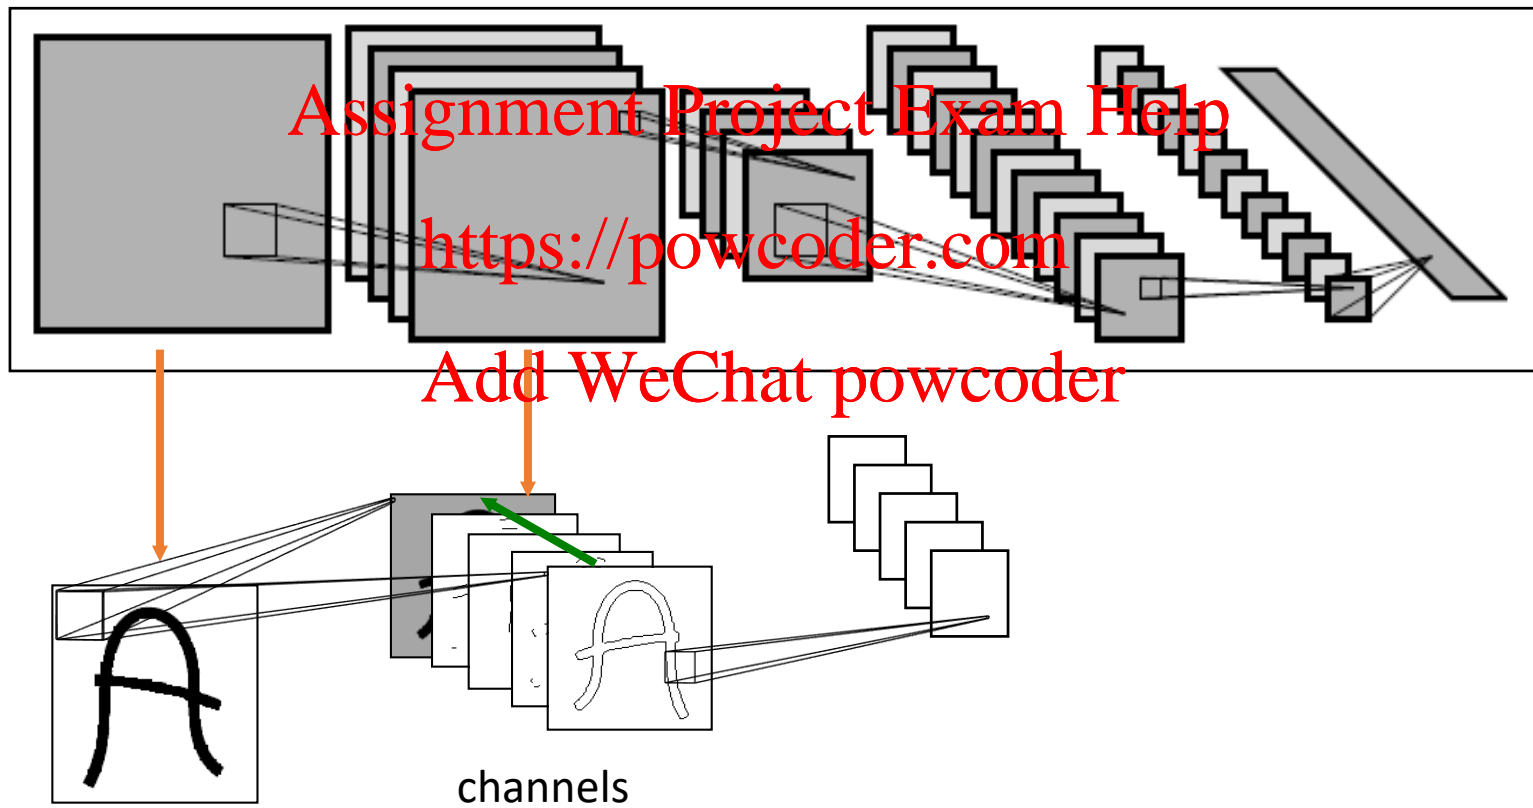


filter

Output

Stacking convolutional layers

- Each layer outputs multi-channel **feature maps** (like images)
- Next layer learns filters on previous layer's feature maps




Pooling layers

- Convolution with stride > 1 reduces the size of the input
- Another way to downsize the feature map is with **pooling**
- A pooling layer subsamples the input in each sub-window
 - **max-pooling**: chose the max in a window
 - **mean-pooling**: take the average

Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder




Pooling layer

- the **pooling** layers reduce the spatial resolution of each feature map
- Goal is to get a **certain degree** of **shift** and **distortion**

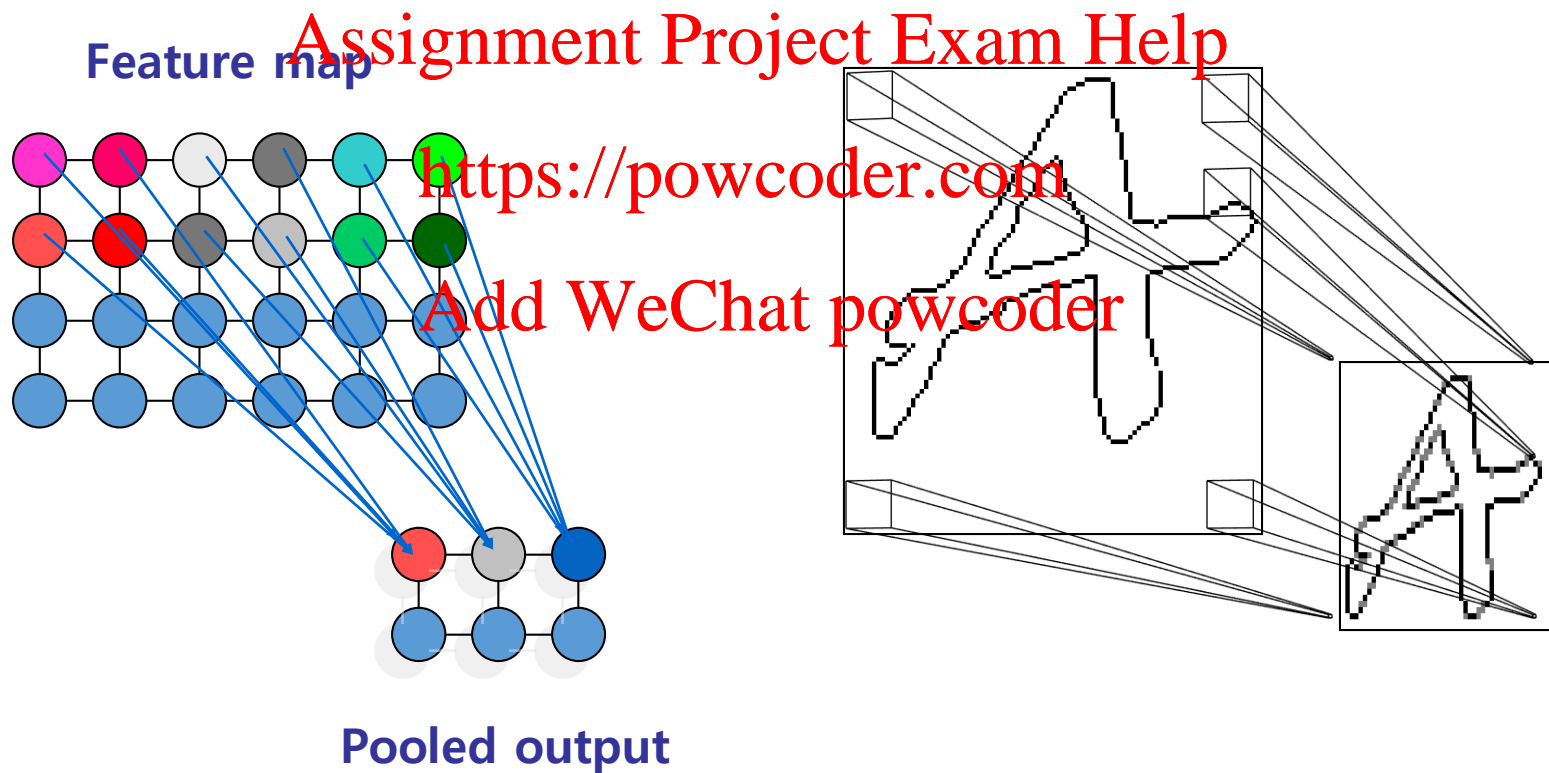


Distortion invariance




Pooling layer

- the **weight sharing** is also applied in pooling layers
- for mean/max pooling, no weights are needed




Putting it all together...



Convolutional Neural Network

A better architecture for 2d signals



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

LeNet

Deep Convolutional Networks

The Unreasonable Effectiveness of Deep Features



Maximal activations of pool_5 units

[R-CNN]

Rich visual structure of features deep in hierarchy.



conv_5 DeConv visualization

[Zeiler-Fergus]



Assignment Project Exam Help

<https://powcoder.com>




Add WeChat powcoder

Convolutional Neural Nets

Why they rule

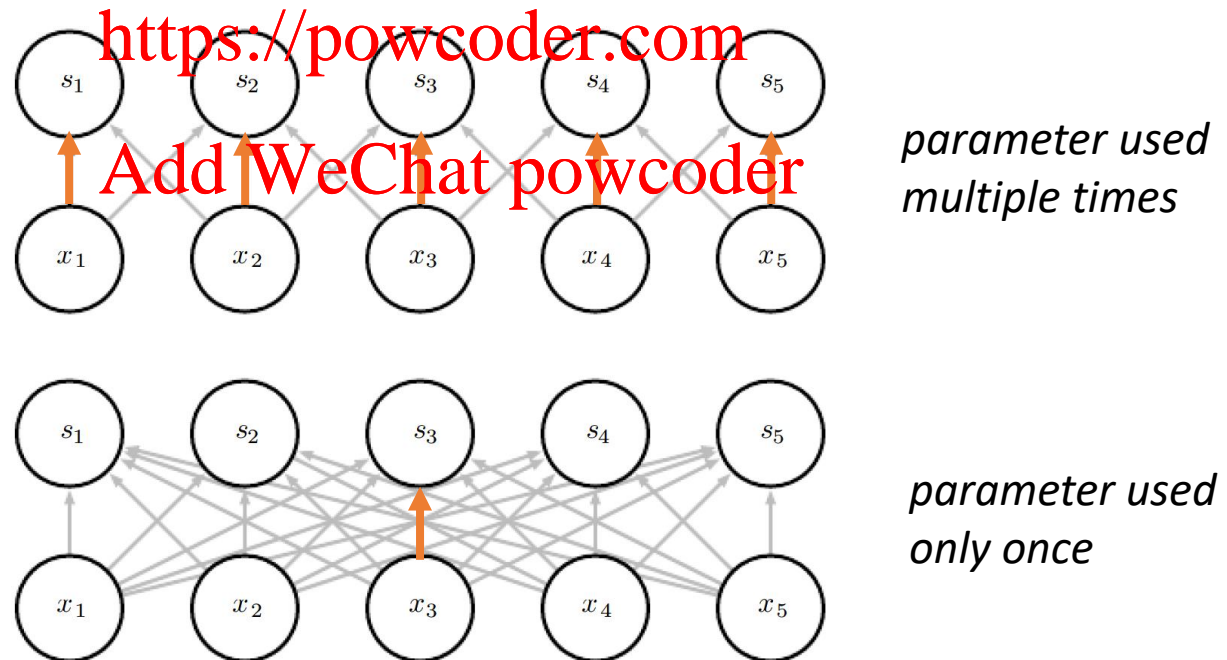
Why CNNs rule: Sparsity

- CNNs have sparse interactions, because the kernel is smaller than the input
- E.g. in thousands or millions pixel image, can detect small meaningful features such as edges
- Very efficient computation!
 - For m inputs and n outputs, matrix multiplication requires $O(m \times n)$ runtime (per example)
 - For k connections to each output, need only $O(k \times n)$ runtime
- Deep layers have larger effective inputs, or **receptive fields**




Why CNNs rule: Parameter sharing

- Kernel weights are shared across all locations
- Statistically efficient – learn from more data
- Memory efficient – store only k parameters, since $k \ll m$, this is much smaller than $m \times n$.



Why CNNs rule: Translation invariance

- Output is invariant to translation of input
 - spatial translation for images
 - temporal translation for time sequences
- useful when some function of a small local window is useful when applied to multiple input locations
- Note, not invariant to other transformations of input, such as large image rotation
- Pooling provides additional invariance to distortions





Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

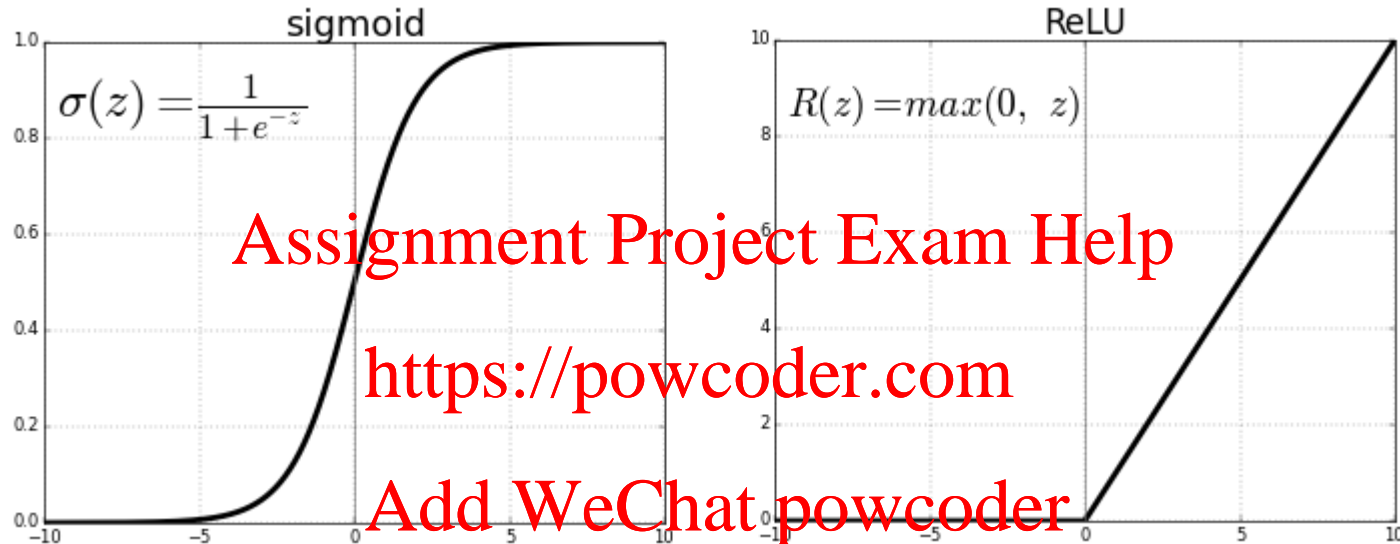
Convolutional Neural Nets

Example


CIFAR-10 Demo ConvJS Network



ReLU: rectified linear unit



ReLU function $g(x) = \max(0, x)$



input (32x32x3)




filter size 5x5x3, stride 1



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




input (32x32x3)



RELU

conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$


filter size 5x5x3, stride 1



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder




input (32x32x3)



filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$




pool (16x16x16)
pooling size 2x2, stride 2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



filter size $5 \times 5 \times 3$, stride 1




input ($32 \times 32 \times 3$)




conv ($32 \times 32 \times 16$) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$

pool ($16 \times 16 \times 16$)
pooling size 2×2 , stride 2



filter size $5 \times 5 \times 16$, stride 1



https://powcoder.com


Add WeChat powcoder

• • •


RELU



conv ($16 \times 16 \times 20$) params: $20 \times 5 \times 5 \times 16 + 20 = 8020$



pool ($8 \times 8 \times 20$)
pooling size 2×2 , stride 2



input (32x32x3)



One more conv+RELU+pool:

conv (8x8x20)

filter size 5x5x20, stride 1

relu (8x8x20)

pool (4x4x20)

pooling size 2x2, stride 2

parameters: $20 \times 5 \times 5 \times 20 + 20 = 10020$

Assignment Project Exam Help

<https://powcoder.com>

softmax (1x1x10)

Add WeChat powcoder

dog

car

cat

:

Testing the network

- Show top three most likely classes



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

Next Class

Neural Networks IV: Recurrent Nets:

recurrent networks; training strategies

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder