

Announcements

Reminder: pset5 self-grading form and pset6 out Thursday, due 11/19 (1 week)

Assignment Project Exam Help

- No lab this week!

<https://powcoder.com>

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Reinforcement Learning II

Recall: MDP notation

- S – set of States
- A – set of Actions
- $R: S \rightarrow \mathbb{R}$ (Reward)
- P_{sa} – transition probabilities ($p(s, a, s') \in \mathbb{R}$)
- γ – discount factor

$$\text{MDP} = (S, A, R, P_{sa}, \gamma)$$



MDP (Simple example)



MDP (Simple example)

- States S = locations
- Actions $A = \{\uparrow, \rightarrow, \leftarrow, \downarrow\}$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

	1	2	3	4	
1					✓
2					✗
3					

MDP (Simple example)

- States S = locations
- Actions $A = \{\uparrow, \rightarrow, \leftarrow, \downarrow\}$
- Reward $R: S \rightarrow \mathbb{R}$
- Transition P_{sa}

	1	2	3	4
1	-0.02	-0.02	-0.02	+1
2	-0.02		-0.02	-1
3	-0.02	-0.02		-0.02

$$P_{(3,3),\uparrow}((2,3)) = 0.8$$

$$P_{(3,3),\uparrow}((3,4)) = 0.1$$

$$P_{(3,3),\uparrow}((3,2)) = 0.1$$

$$P_{(3,3),\uparrow}((1,3)) = 0$$

$$\vdots$$

MDP - Dynamics

- Start from state S_0
- Choose action A_0
- Transit to $S_1 \sim P_{S_0 A_0}$

	1	2	3	4
1	-.02	-.02	-.02	+1
2	-.02		-.02	-1
3				-.02

- Continue... <https://powcoder.com>

-.02

-.02

-.02

Add WeChat powcoder

- Total payoff:

-.02

-.02

-.02

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Q-Learning (discrete)

Reinforcement Learning

Q-value function

- ▶ A **value function** is a prediction of future reward
 - ▶ “How much reward will I get from action a in state s ?”
- ▶ **Q**-value function gives expected total reward
 - ▶ from state s and action a
 - ▶ under policy π
 - ▶ with discount factor γ

$$Q^\pi(s, a) = \mathbb{E} [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \mid s, a]$$

- ▶ Value functions decompose into a Bellman equation

$$Q^\pi(s, a) = \mathbb{E}_{s', a'} [r + \gamma Q^\pi(s', a') \mid s, a]$$

Optimal Q-value function

- ▶ An optimal value function is the maximum achievable value

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- ▶ Once we have Q^* , we can act optimally.

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- ▶ Optimal value maximises over all decisions. Informally:

$$\begin{aligned} Q^*(s, a) &= r_{t+1} + \gamma \max_{a_{t+1}} r_{t+2} + \gamma^2 \max_{a_{t+2}} r_{t+3} + \dots \\ &= r_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \end{aligned}$$

- ▶ Formally, optimal values decompose into a Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

Q-learning algorithm

The agent interacts with the environment, updates Q recursively

Assignment, Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

```
initialize Q-table, run_actions, arbitrarily  
observe initial state s  
repeat  
    select and carry out an action a  
    observe reward r and new state s'  
     $Q[s,a] = Q[s,a] + \alpha(r + \gamma \max_{a'} Q[s',a'] - Q[s,a])$   
    s = s'  
until terminated
```

current value

learning rate

discount

largest increase over all
possible actions in new state

Q-learning example






Goal: get from bottom left to top right



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



<https://www.youtube.com/watch?v=R88CiN7dTZc>

Exploration vs exploitation

- How does the agent select actions during learning? Should it trust the learned values of $Q(s, a)$ to select actions based on it? or try other actions hoping this may give it a better reward?
- This is known as the exploration vs exploitation dilemma
- Simple ϵ -greedy approach: at each step with small probability ϵ , the agent will pick a random action (explore) or with probability $(1-\epsilon)$ the agent will select an action according to the current estimate of Q-values
- The ϵ value can be decreased overtime as the agent becomes more confident with its estimate of Q-values



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Continuous state

Reinforcement Learning

Continuous state - Pong


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

<https://www.youtube.com/watch?v=YOW8m2YGtRg>

MDP for Pong



In this case, what are these?

- S – set of States
- A – set of Actions
- $R: S \rightarrow \mathbb{R}$ (Reward)
- P_{sa} – transition probabilities ($p(s, a, s') \in \mathbb{R}$)

<https://powcoder.com>
Add WeChat powcoder

Can we learn Q-value?

- Can discretize state space, but it may be too large
- Can simplify state by adding domain knowledge (e.g. paddle, ball), but it may not be available
- Instead, use a neural net to learn good features of the state!



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Deep RL

Reinforcement Learning

Deep RL playing DOTA



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

https://www.youtube.com/watch?v=eHipy_j29Xw

Deep RL

- V , Q or π can be approximated with deep network

- Deep Q-Learning

- Input: state, action
 - Output: Q-value

- Alternative: learn a Policy Network

- Input: state
 - Output: distribution over actions

Assignment Project Exam Help

<https://powcoder.com>

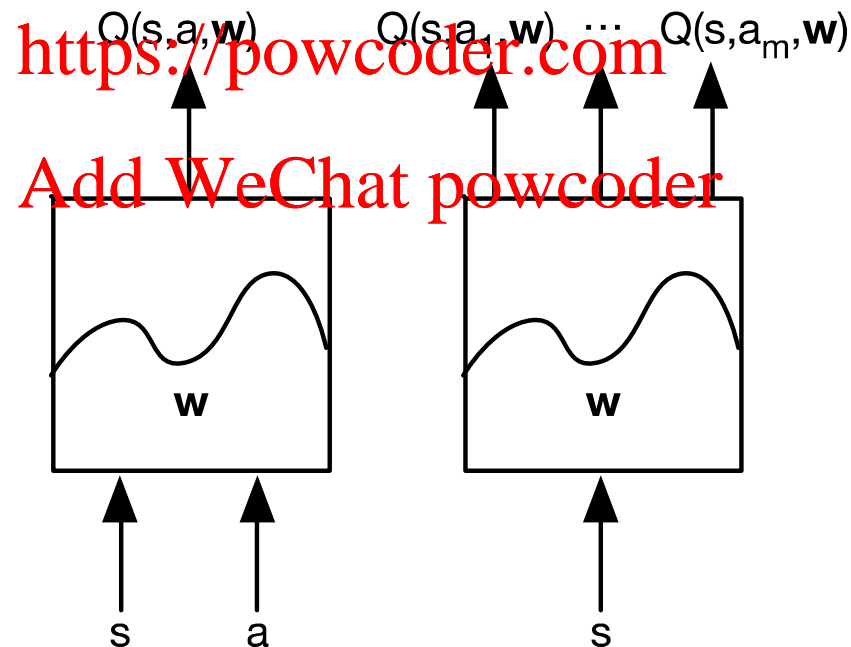
Add WeChat powcoder

Q-value network

Represent value function by **Q-network** with weights **w**

$$Q(s, a, \mathbf{w}) \approx Q^*(s, a)$$

Assignment Project Exam Help



Q-value network

- ▶ Optimal Q-values should obey Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q(s', a')^* \mid s, a \right]$$

- ▶ Treat right-hand side $r + \gamma \max_{a'} Q(s', a', \mathbf{w})$ as a target
- ▶ Minimise MSE loss by stochastic gradient descent

$$l = \left(r + \gamma \max_a Q(s', a', \mathbf{w}) - Q(s, a, \mathbf{w}) \right)^2$$

- ▶ Converges to Q^* using table lookup representation
- ▶ But **diverges** using neural networks due to:
 - ▶ Correlations between samples
 - ▶ Non-stationary targets

Deep Q-network (DQN)

To remove correlations, build data-set from agent's own experience

Assignment Project Exam Help
<https://powcoder.com>
 Add WeChat powcoder

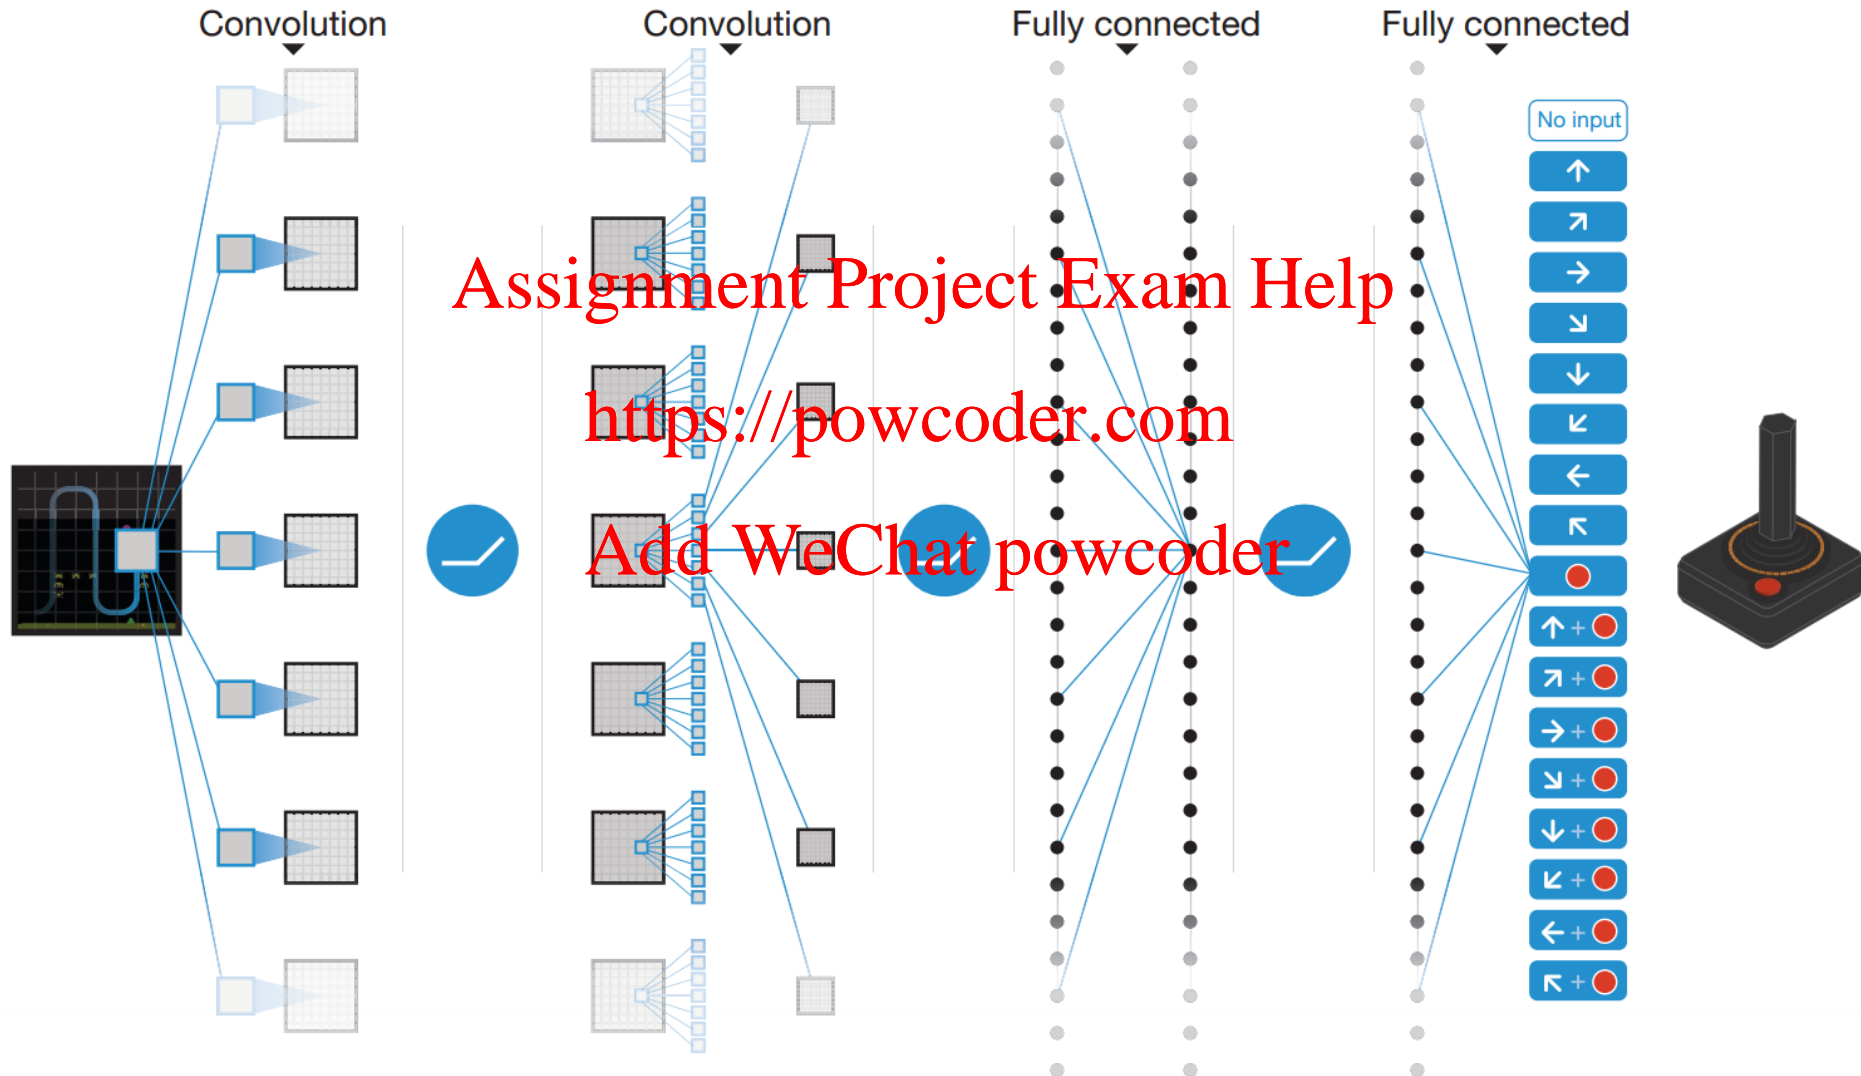
	s_1, a_1, r_2, s_2
	s_2, a_2, r_3, s_3
	s_3, a_3, r_4, s_4
$s_t, a_t, r_{t+1}, s_{t+1}$	$\rightarrow s_t, a_t, r_{t+1}, s_{t+1}$

Sample experiences from data-set and apply update

$$l = \left(r + \gamma \max_{a'} Q(s', a', \mathbf{w}^-) - Q(s, a, \mathbf{w}) \right)^2$$

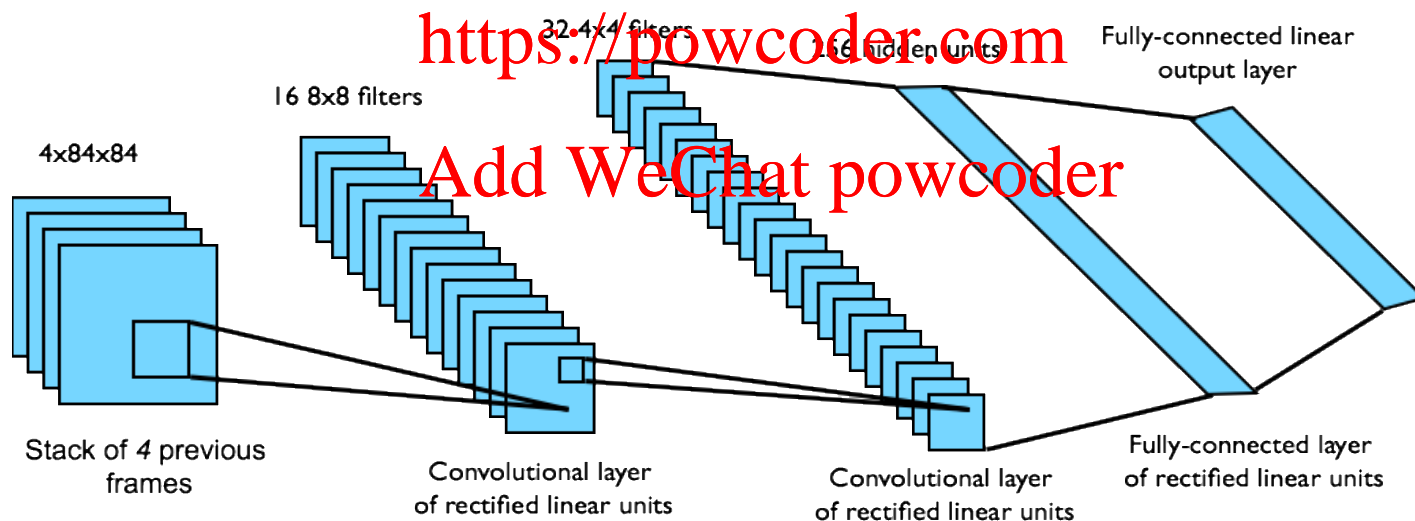
To deal with non-stationarity, target parameters \mathbf{w}^- are held fixed

DQN - Playing Atari



DQN - Playing Atari

- | End-to-end learning of values $Q(s, a)$ from pixels s
- | Input state s is stack of raw pixels from last 4 frames
- | Output is $Q(s, a)$ for 18 joystick/ button positions
- | Reward is change in score for that step



Network architecture and hyperparameters fixed across all games

DQN - Playing Atari

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}


 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

DQN for Atari

DQN paper:

www.nature.com/articles/nature14236

DQN demo:

<https://www.youtube.com/watch?v=qXKQf2BOSE>


DQN source code:

www.sites.google.com/a/deepmind.com/dqn/

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



Deep RL

- V , Q or π can be approximated with deep network
- Deep Q-Learning
 - Input: state, action
 - Output: Q-value
- Alternative: learn a Policy Network
 - Input: state
 - Output: distribution over actions

Policy network for pong




- define a *policy network* that implements the player
- takes the state of the game and decides what to do (move UP or DOWN)
- 2-layer neural network that takes the raw image pixels* (100,800 = 210x160x3), outputs the probability of going UP

*feed at least 2 frames to the policy network so that it can detect motion.

<http://karpathy.github.io/2016/05/31/rl/>

Policy gradient



- Suppose network predicts

$$p(\text{UP}) = 30\%$$

$$p(\text{DOWN}) = 70\%$$

- Can sample an action from this distribution and execute it
- Can immediately use gradient of 1.0 for DOWN and backprop to find the gradient vector that would encourage the network to predict DOWN

Assignment Project Exam Help

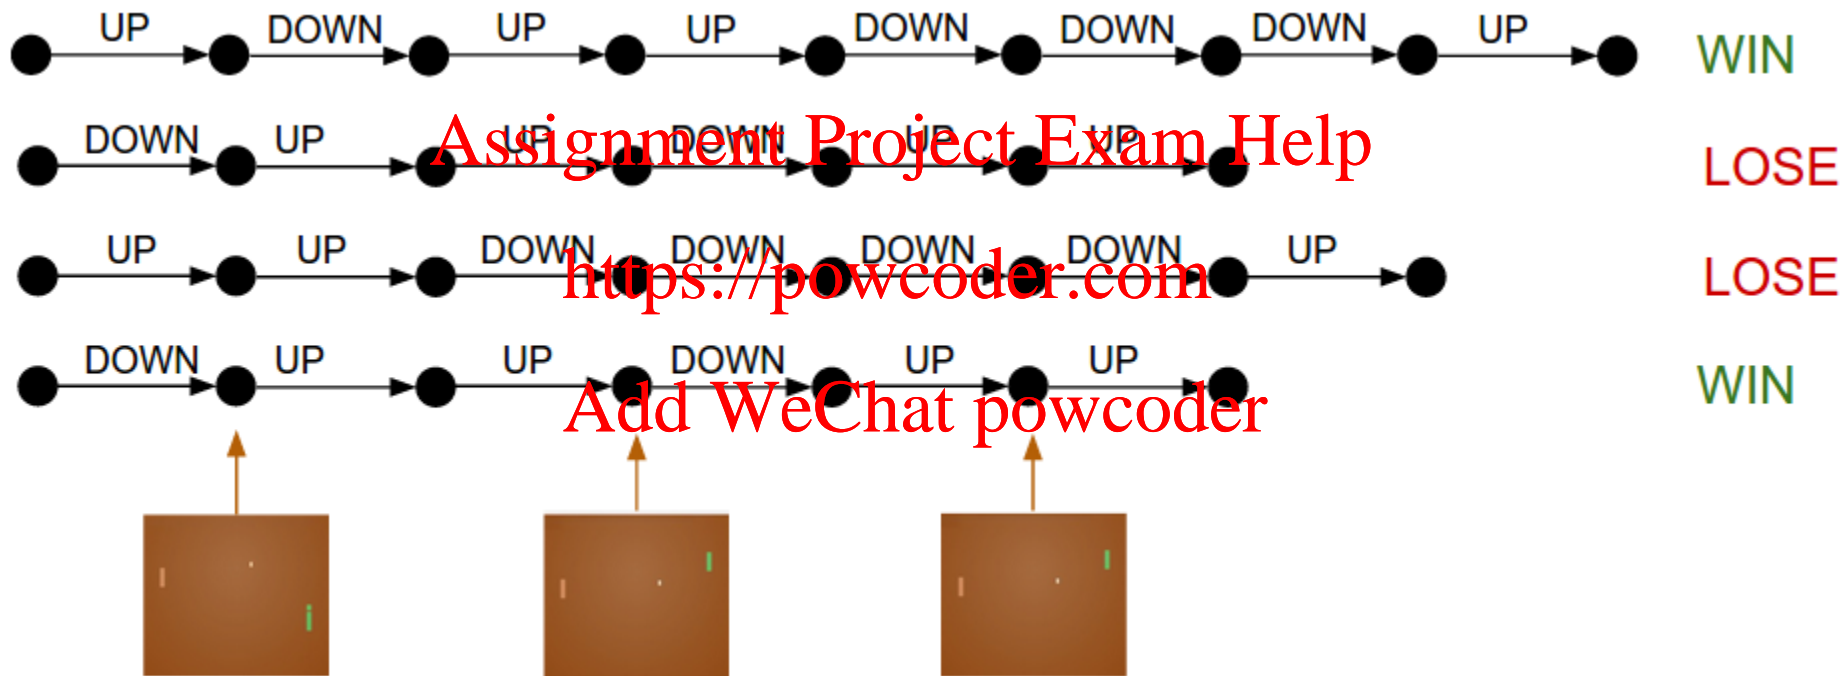
<https://powcoder.com>

Add WeChat powcoder

Problem: do not yet know if going DOWN is good!

Solution: simply wait until the end of the game, then take the reward we get (either +1 if we won or -1 if we lost), and enter that as the gradient for taken actions

Policy gradient



Problems with this?

- what if we made a good action in frame 50 (bouncing the ball back correctly), but then missed the ball in frame 150?
- If every single action is now labeled as bad (because we lost), wouldn't that discourage the correct bounce on frame 50?
- Yes, but after thousands/millions of games, network will learn a good policy

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Policy gradient

Want to maximize

$$E_{x \sim p(x|\theta)}[f(x)]$$

<https://powcoder.com>

$f(x)$ is the reward function

Add WeChat powcoder

$p(x)$ is the policy network with parameters θ

(i.e. change the network's parameters so that action samples get higher rewards)

Downsides of RL

- RL is less sampling efficient than supervised learning because it involves bootstrapping, which uses an estimate of the Q-value to update the Q-value predictor
- Rewards are usually sparse and learning requires to reach the goal by chance
- Therefore, RL might not find a solution at all if the state space is large or if the task is difficult

References

Andrew Ng's Reinforcement Learning course, lecture 16

<https://www.youtube.com/watch?v=Rtxl449ZjSc>

Assignment Project Exam Help

Andrej Karpathy's blog post on policy gradient

<http://karpathy.github.io/2016/05/31/rl/>

<https://powcoder.com>

Mnih et. al, Playing Atari with Deep Reinforcement Learning (DeepMind)

<https://www.cs.toronto.edu/~vmnih/docs/playing-atari.pdf>

Add WeChat powcoder

Intuitive explanation of deep Q-learning

<https://www.nervanasys.com/demystifying-deep-reinforcement-learning/>

Next Class

Unsupervised Learning III: Anomaly Detection

Anomaly detection methods: density estimation, reconstruction based method, One Class SVM; evaluating anomaly detection

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder